Microsoft

Microsoft Cloud Workshop

Big data and visualization

Julius Chen

Abstract and learning objectives

Abstract

Deploy a web app using Machine Learning to predict travel delays given flight delay data and weather conditions. Plan a bulk data import operation, followed by preparation, such as cleaning and manipulating the data for testing, and training your Machine Learning model.

Learning objectives

- Build a complete Azure Machine Learning (ML) model.
- Integrate an Azure ML web service into a Web App.
- Use Azure Data Factory (ADF) for data movement and operationalizing ML scoring.
- Summarize data with HDInsight and Spark SQL.
- Visualize batch predictions on a map using Power BI.

Step 1: Review the customer case study

Outcome

Analyze your customer needs

Timeframe

15 minutes

Customer situation

AdventureWorks Travel (AWT) provides concierge services for business travelers.

Interested in using predictive analytics to differentiate themselves in an increasingly crowded market.

Customer situation

- Proposed solution to provide flight delay risk assessment to customers
- Plan to use 30 years of flight delay and weather data
- Want to pilot the solution internally

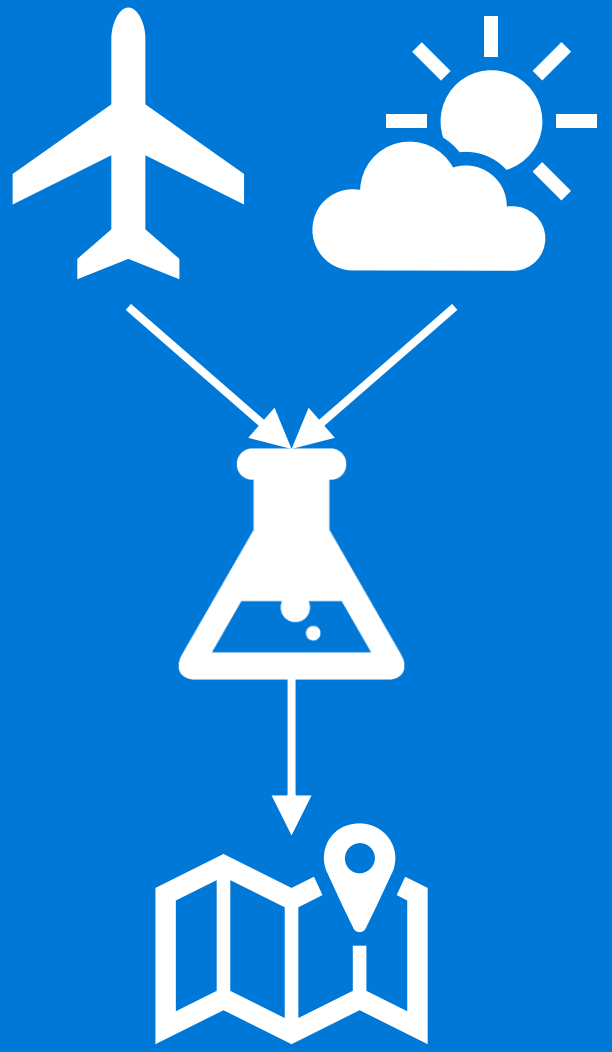

Customer needs

- Modernize their analytics platform
- Ability to query data using SQL
- Load and store all data in Azure
- Use current weather forecast for flight delay predictions
- Proof of concept machine learning model
- Web-based visualizations of flight delay predictions

Customer objections

- Does Azure Machine Learning require a PhD in statistics?

- How long does it take to create and operationalize a machine learning model?

- Can operationalized ML models be flexible in the inputs they support?

Customer objections

- What are the options for running SQL on Hadoop solutions in Azure?

- Does Azure offer anything to speed up querying files in HDFS?

- How can we identify, monitor, and protect PII data?

Customer objections

- Is Azure Data Lake a good fit for our PoC?

- Can access to our SQL DW be limited using Azure Active Directory?

- What data visualization tools are available on Azure? Can access to these be managed with Active Directory?

Common scenarios

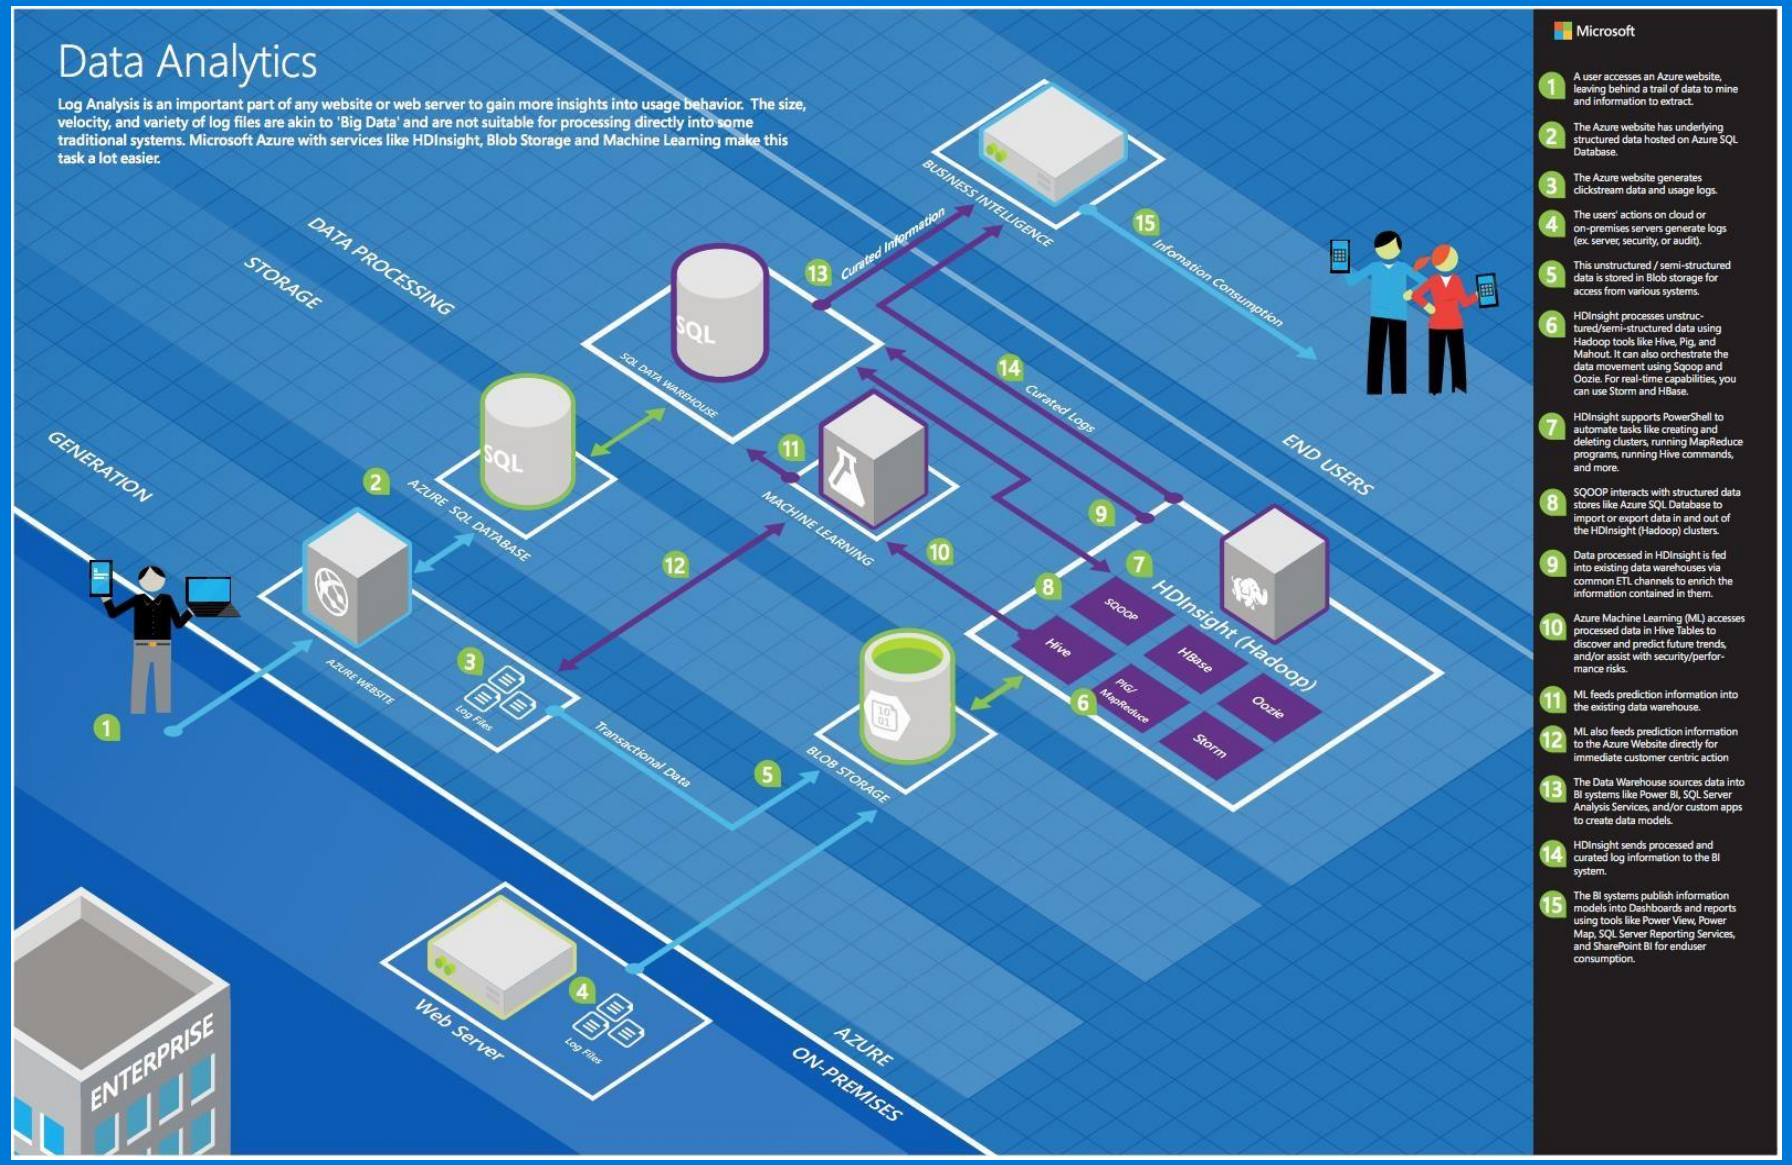

Step 2: Design the solution

Outcome

Design a solution and prepare to present the solution to the target customer audience in a 10-minute chalk-talk format.

Timeframe

60 minutes

Business needs (10 minutes)	• Respond to questions outlined in your guide and list the answers on a flipchart.
Design (35 minutes)	• Design a solution for as many of the stated requirements as time allows. Show the solution on a flipchart.
Prepare (15 minutes)	• Identify any customer needs that are not addressed with the proposed solution. • Identify the benefits of your solution. • Determine how you will respond to the customer's objections. • Prepare for a 10-minute presentation to the customer.

Step 3: Present the solution

Outcome

Present a solution to the target customer in a 10-minute chalk-talk format

Timeframe

30 minutes (15 minutes for each team to present and receive feedback)

Directions

- Pair with another table
- One table is the Microsoft team and the other table is the customer
- The Microsoft team presents their proposed solution to the customer
- The customer asks one of the objections from the list of objections in the case study
- The Microsoft team responds to the objection
- The customer team gives feedback to the Microsoft team

Wrap-up

Outcome

Identify the preferred solution for the case study

Identify solutions designed by other teams

Timeframe

15 minutes

Preferred target audience

- Jack Tradewinds, CIO of AdventureWorks Travel

- The primary audience is business decision makers and technology decision makers.

- Usually we talk to the Infrastructure Managers who report to the CIOs, or to application sponsors (like a VP LOB, CMO) or to those that represent the Business Unit IT or developers that report to application sponsors.

Preferred solution

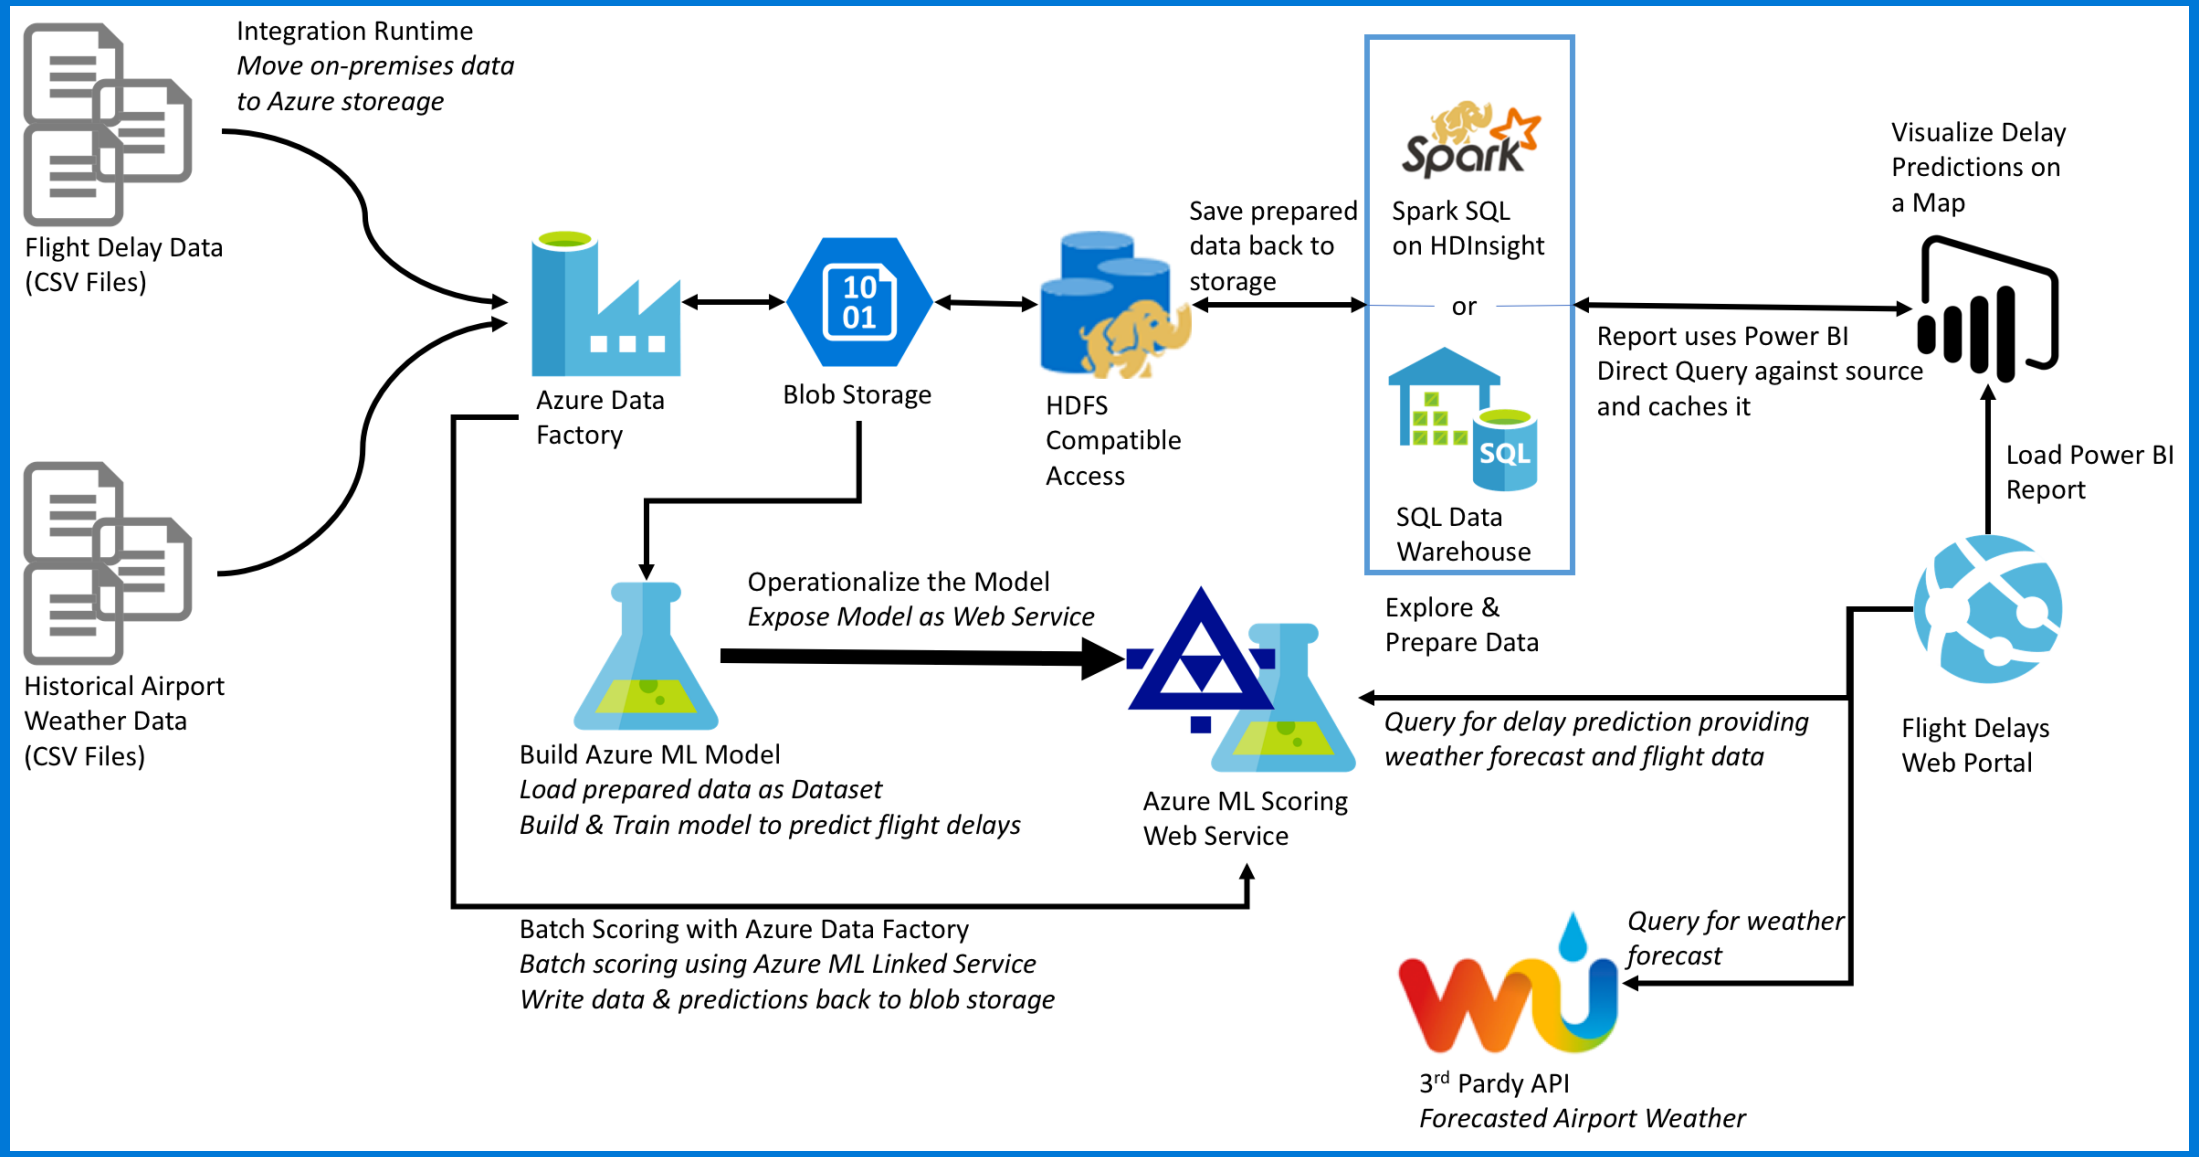

Preferred solution

Data Loading

Historical flight
& weather data

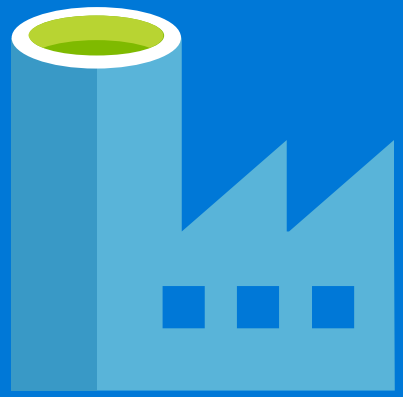

Azure data factory

Azure blob storage

Preferred solution

Data reparation

Spark cluster
on HDInsight

Explore & prepare data

Spark SQL

Jupyter notebook used
by AWT analysts

Preferred solution

Machine learning model

Spark cluster
on HDInsight

Prep data with Spark SQL

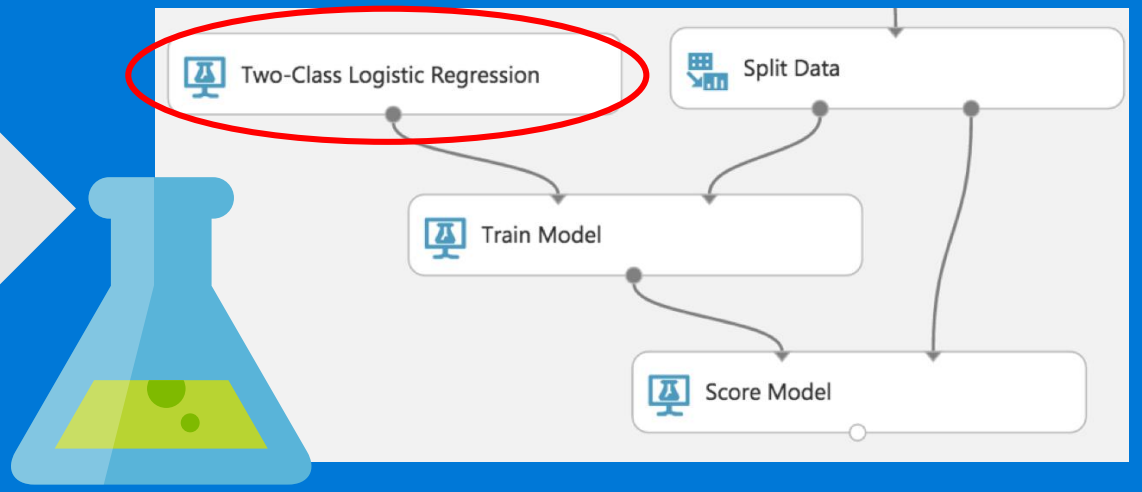

Azure Machine Learning (ML) with
Two-class logistic regression

Preferred solution

Machine learning model

- Start with domain knowledge
- Remove fields that do not add value
- Validate preliminary model against training data

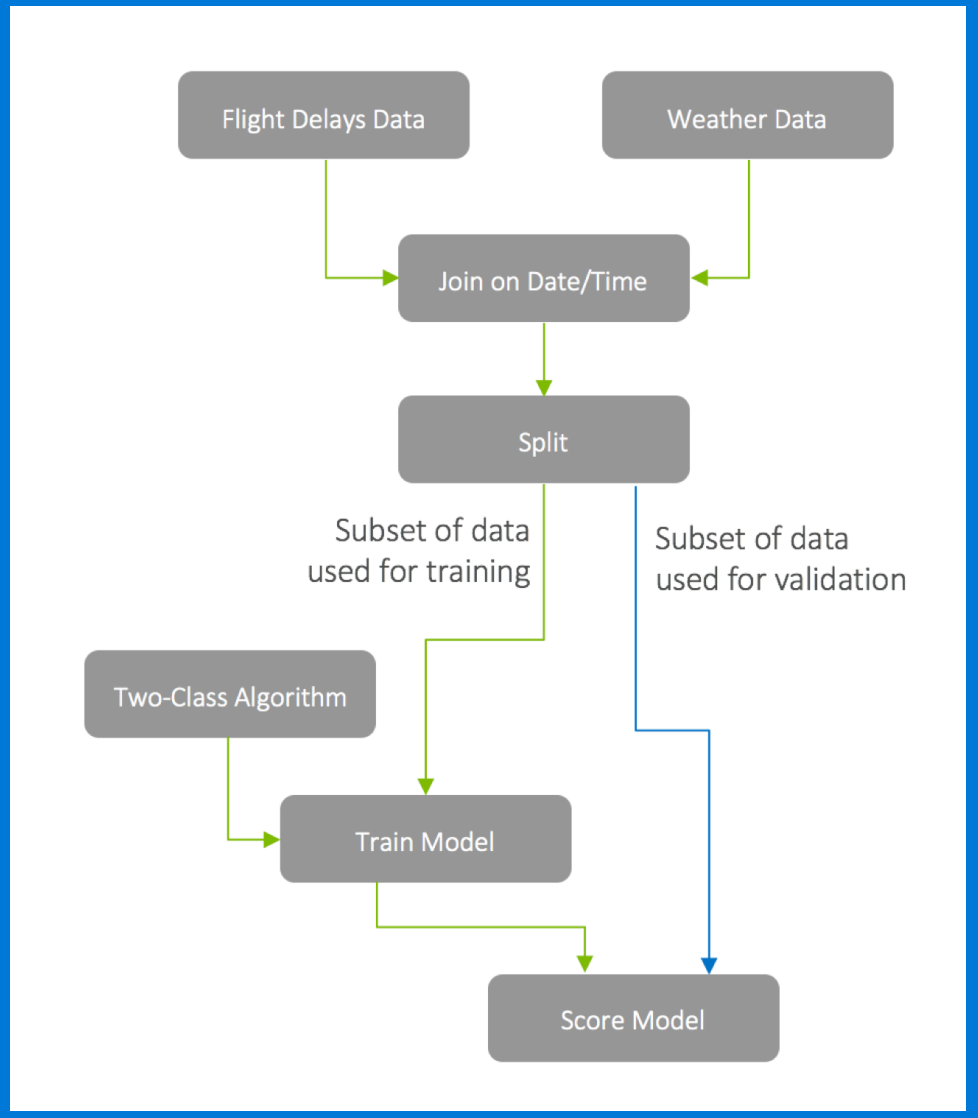

Preferred solution

Machine learning model

- Data lunging with R or Python

- Reserve some historical data to “test” the model

- Measure error on the training set and validation sets separately for indicator of whether model is in danger overfitting.

Preferred solution

Operationalizing machine learning

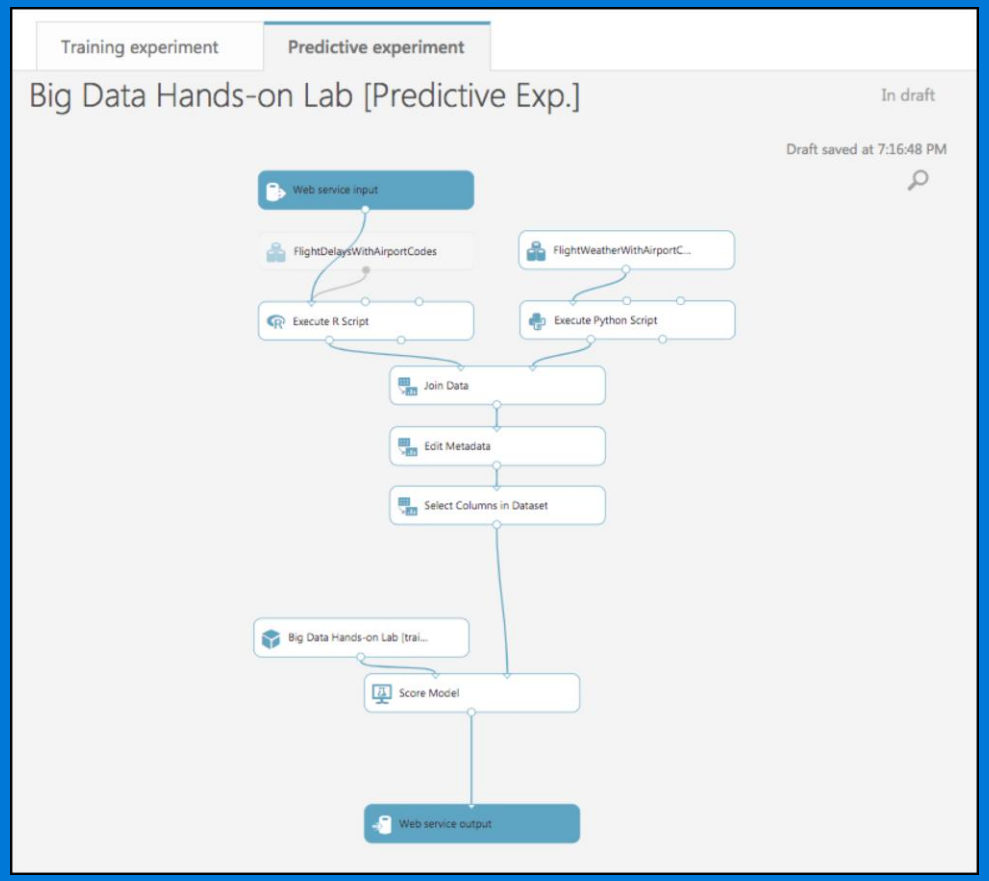

Azure ML Studio

Publish via Azure ML Studio

Operationalize

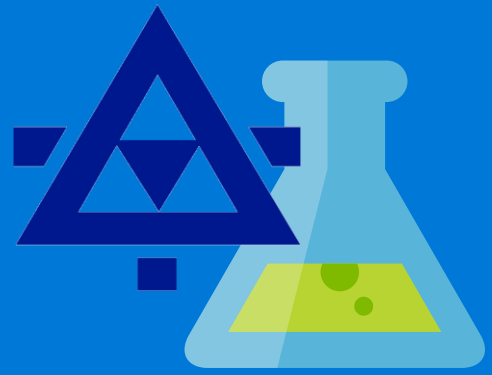

Predictive web service (REST API)

Preferred solution

Operationalizing machine learning

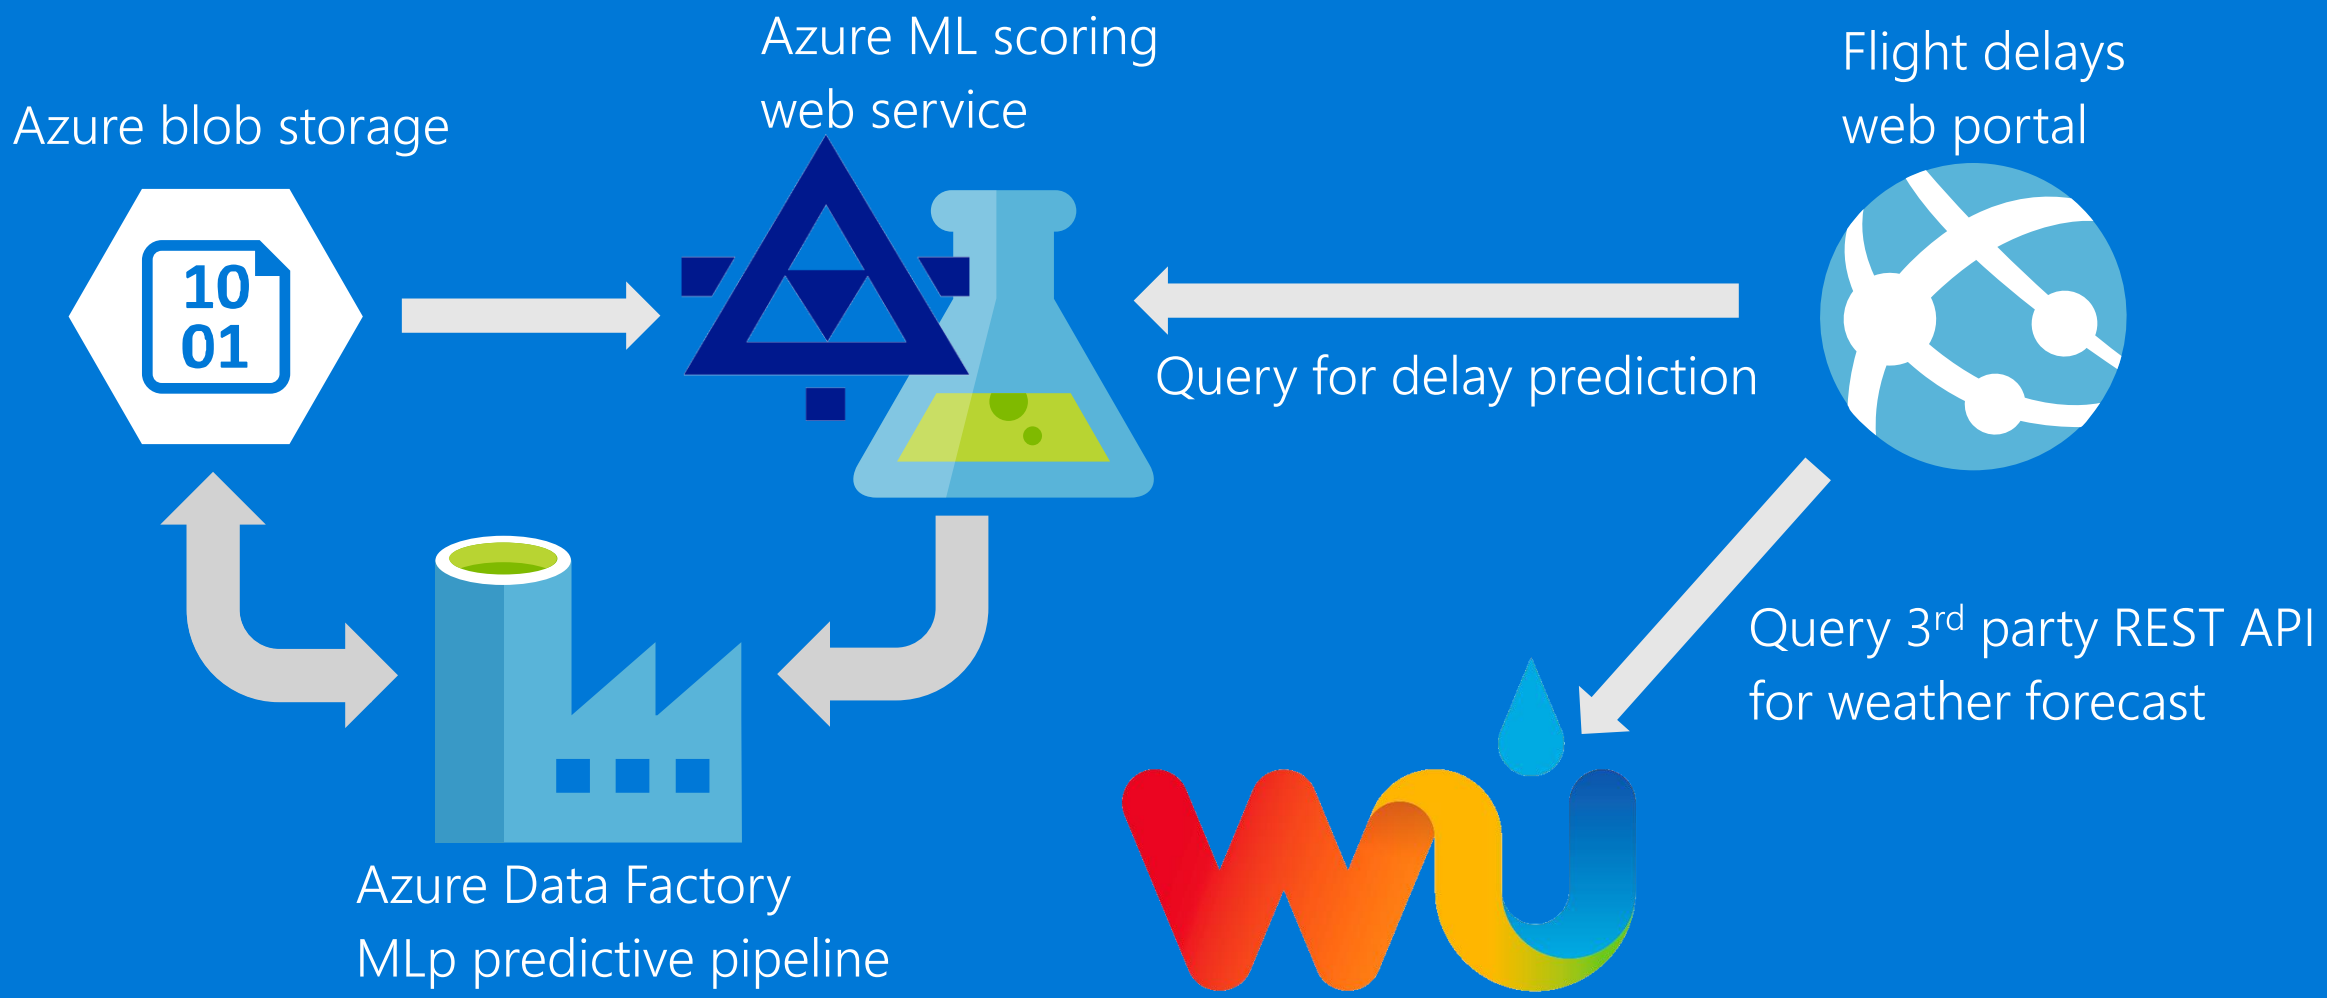

Preferred solution

Visualization and reporting

- Power BI is a good option
- Direct Query against Spark Hive tables.
- Use map visualization

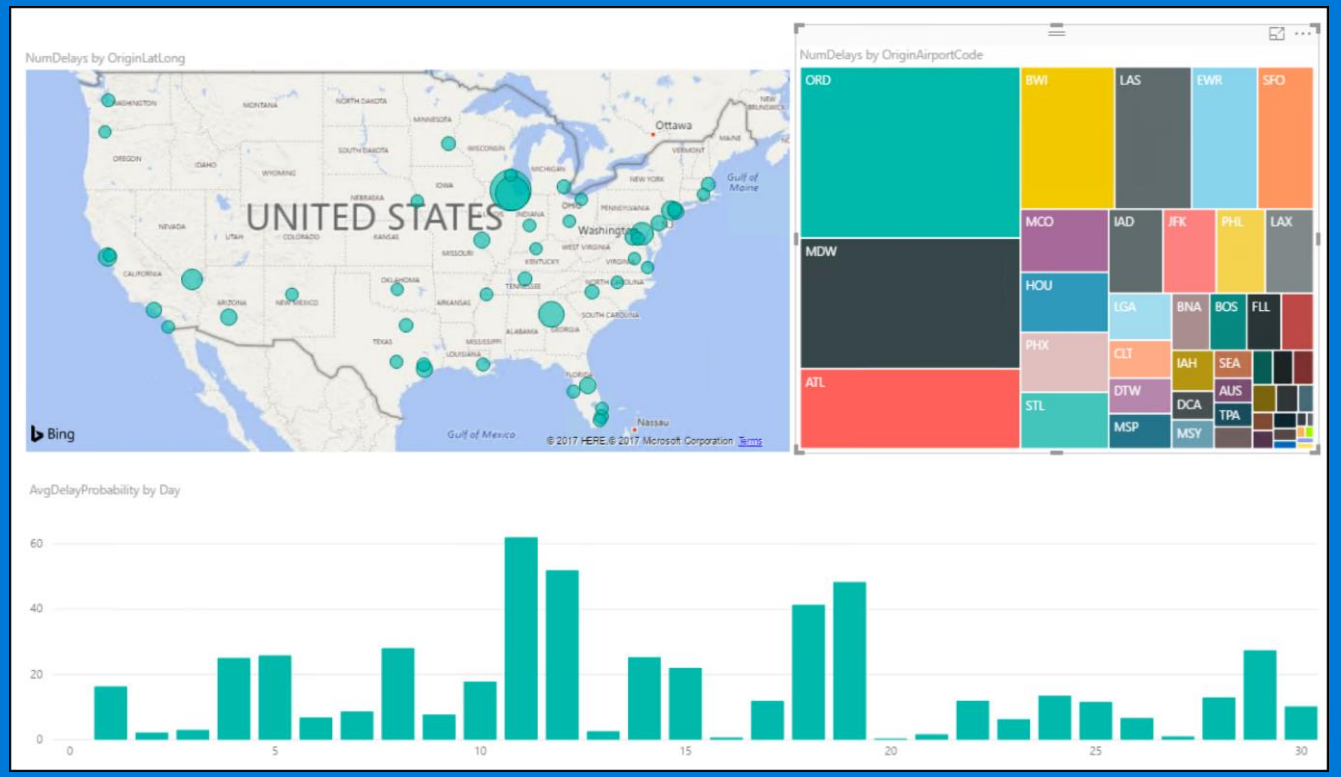

Preferred solution

Visualization and reporting

- Use Query Editor component of the Power BI Desktop, then upload to Power BI service.

- Create content pack with Power BI

- Restrict access in Azure AD

Customer objections

- Does Azure Machine Learning require a PhD in statistics?

- How long does it take to create and operationalize a machine learning model?

- Can operationalized ML models be flexible in the inputs they support?

Customer objections

- What are the options for running SQL on Hadoop solutions in Azure?

- Does Azure offer anything to speed up querying files in HDFS?

- How can we identify, monitor, and protect PII data?

Customer objections

- Is Azure Data Lake a good fit for our PoC?

- Can access to our SQL DW be limited using Azure Active Directory?

- What data visualization tools are available on Azure? Can access to these be managed with Active Directory?

Customer quote

"We are flying into the future with Azure, helping our customers more aggressively schedule their travel, and optimize their non-travel time."

- Jack Tradewinds, CIO of AdventureWorks Travel

Azure Data Services

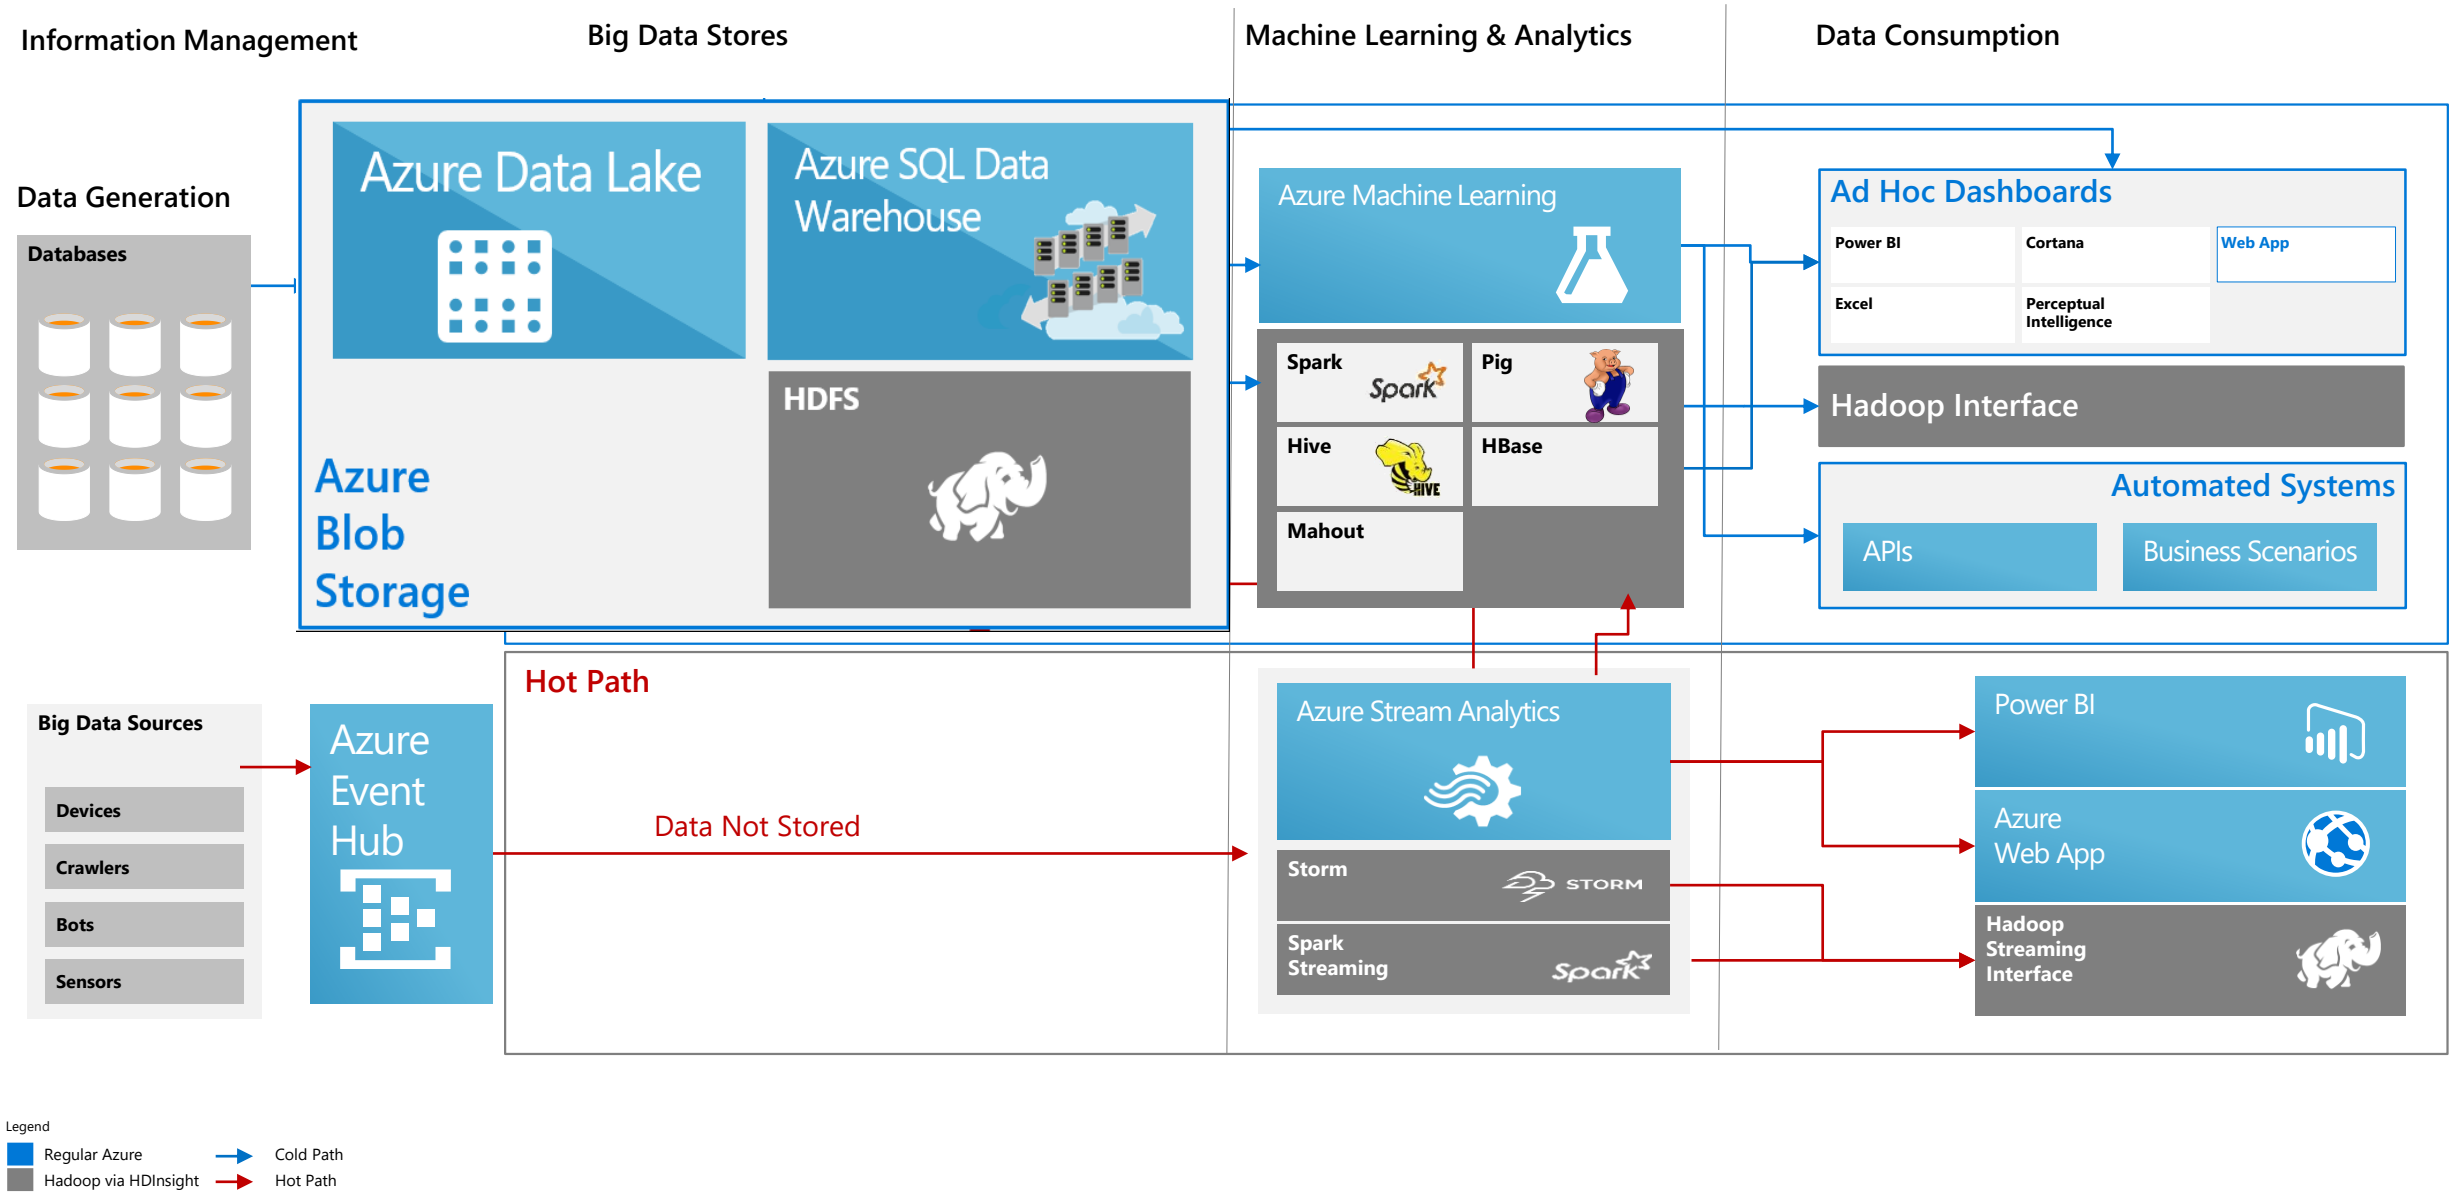

Legend

- Regular Azure
- Hadoop via HDInsight
- Cold Path
- Hot Path

Solution Architecture

Cloud Product | IaaS | New Services | On-Premise

Data Onboarding
- Networking
- Protocol
- Update Approach

Data Storage
- On the fly
- Schema-less
- Schema-rich

Data Processing
- Transformation and Hybrid Data Warehousing
- Machine Learning
- Query Approach

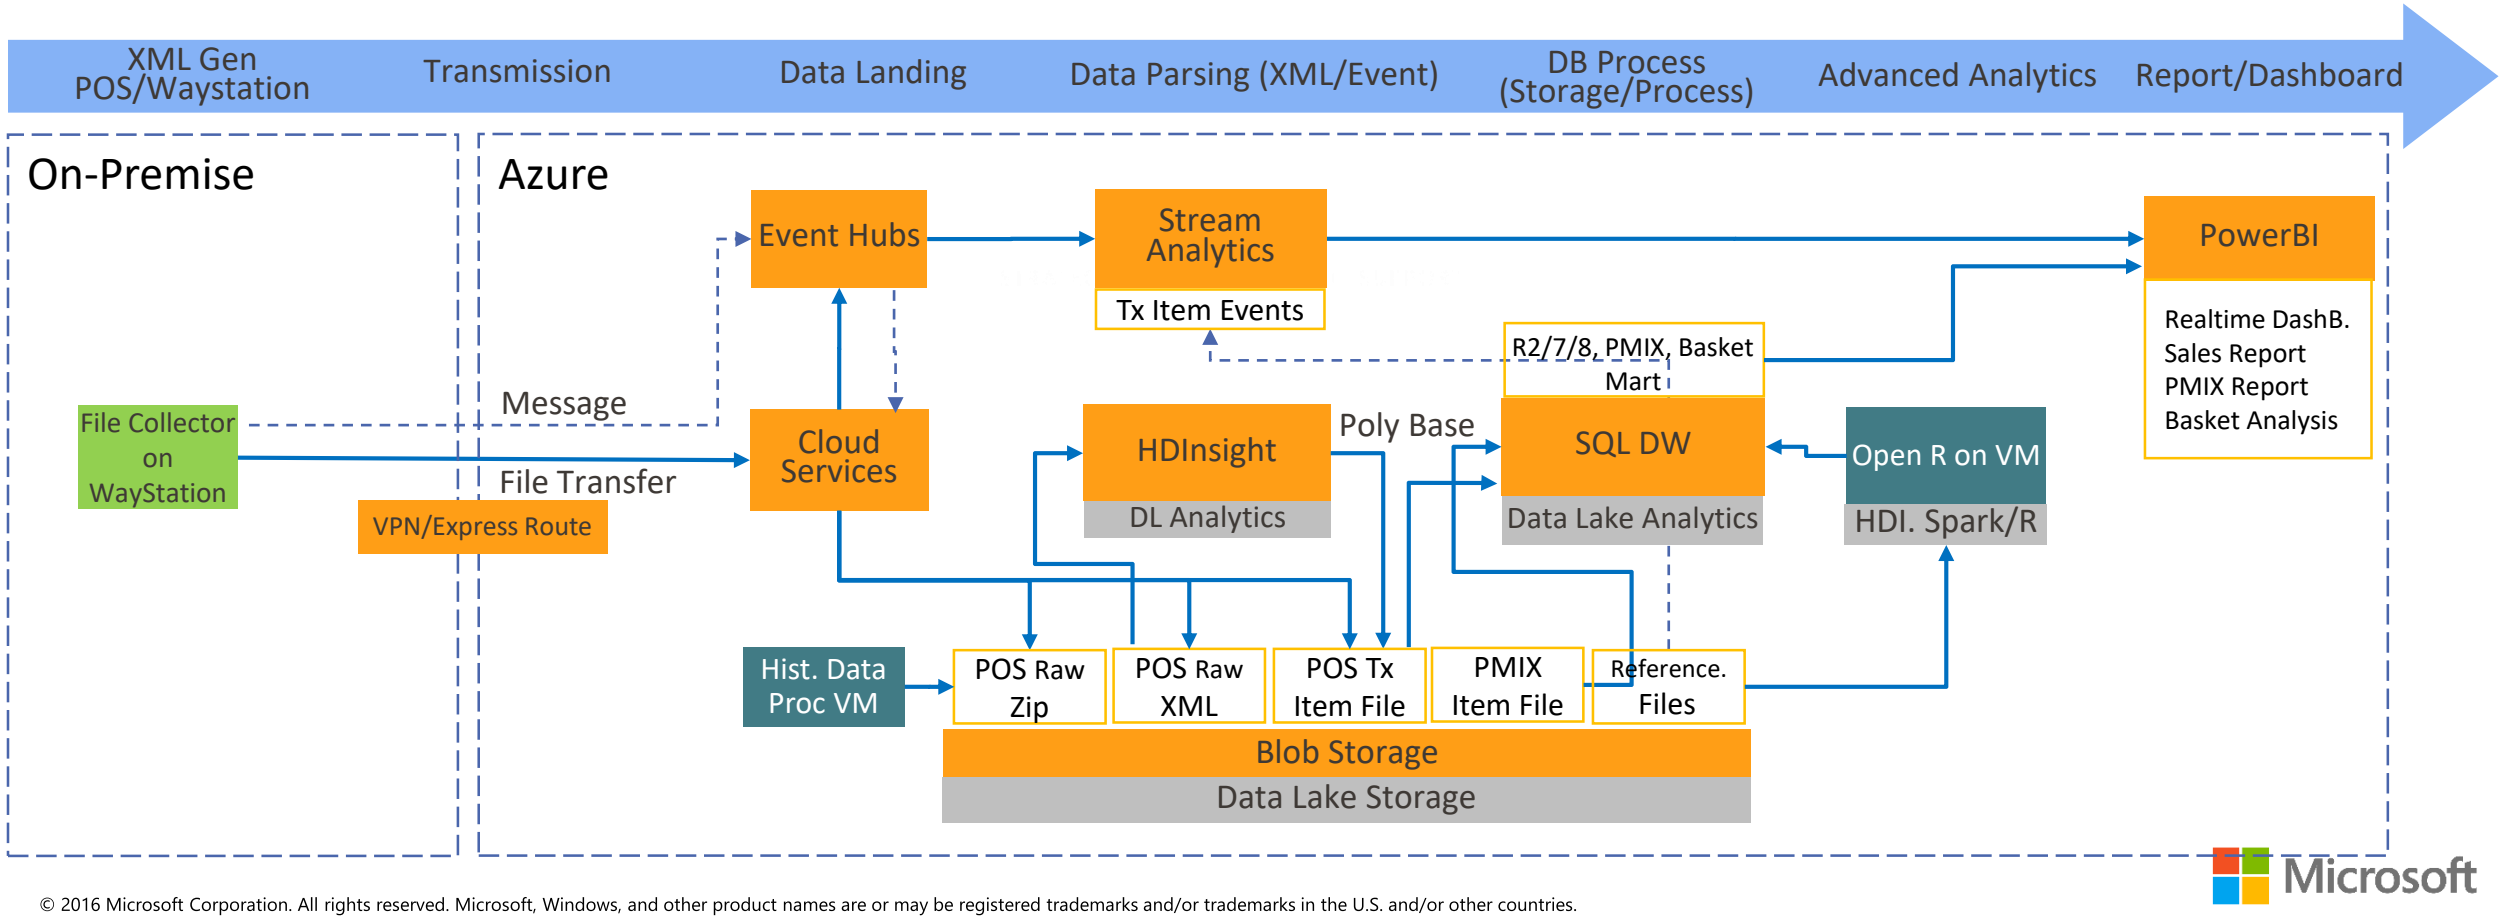

Microsoft

Customer with SQL Server

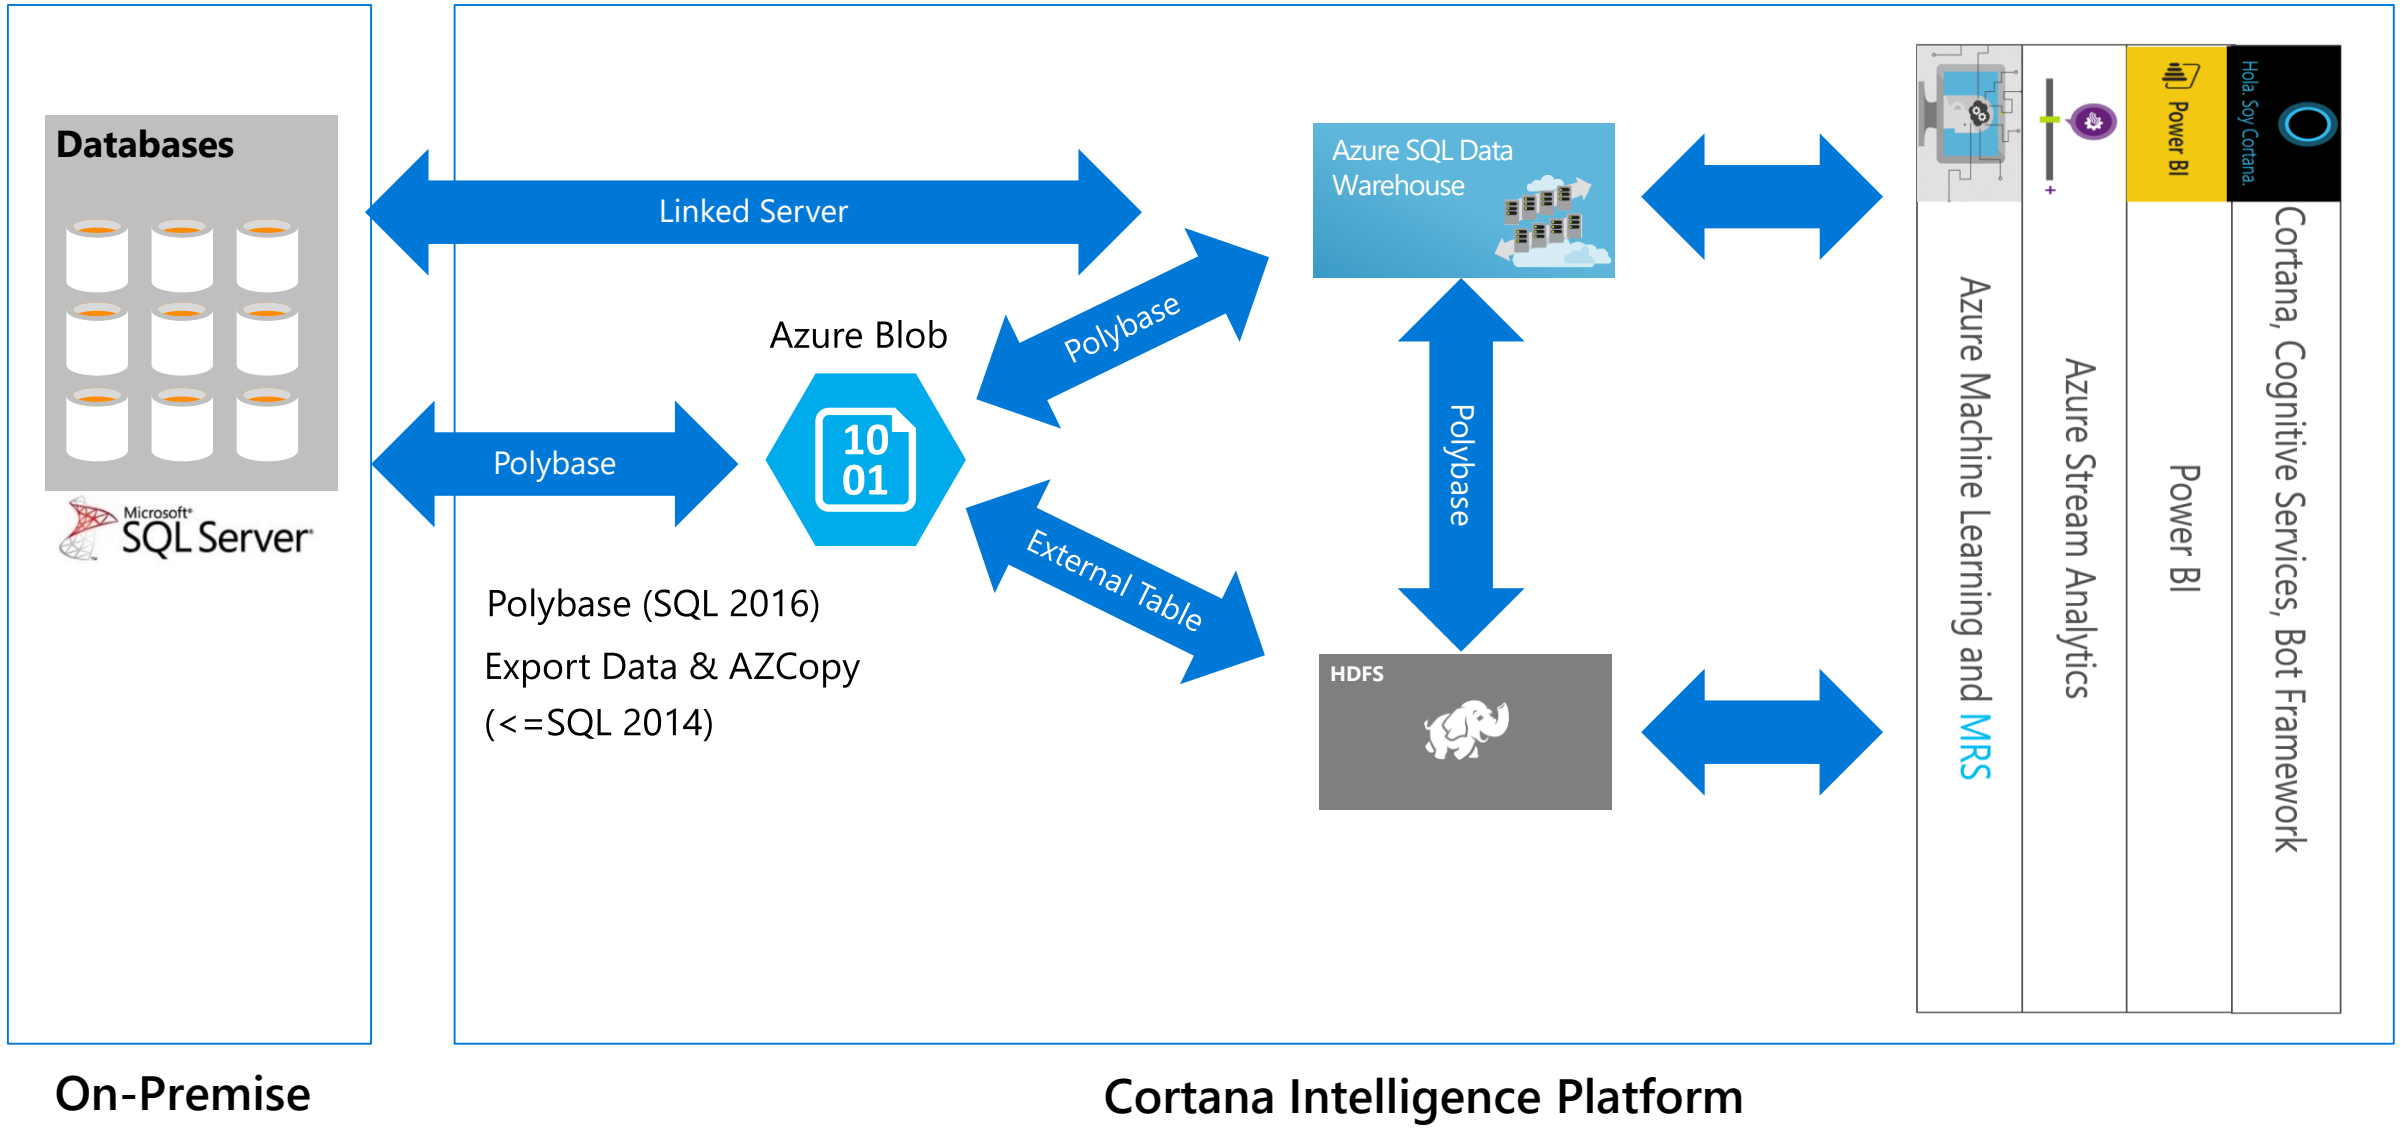

Customer with "Other" DB

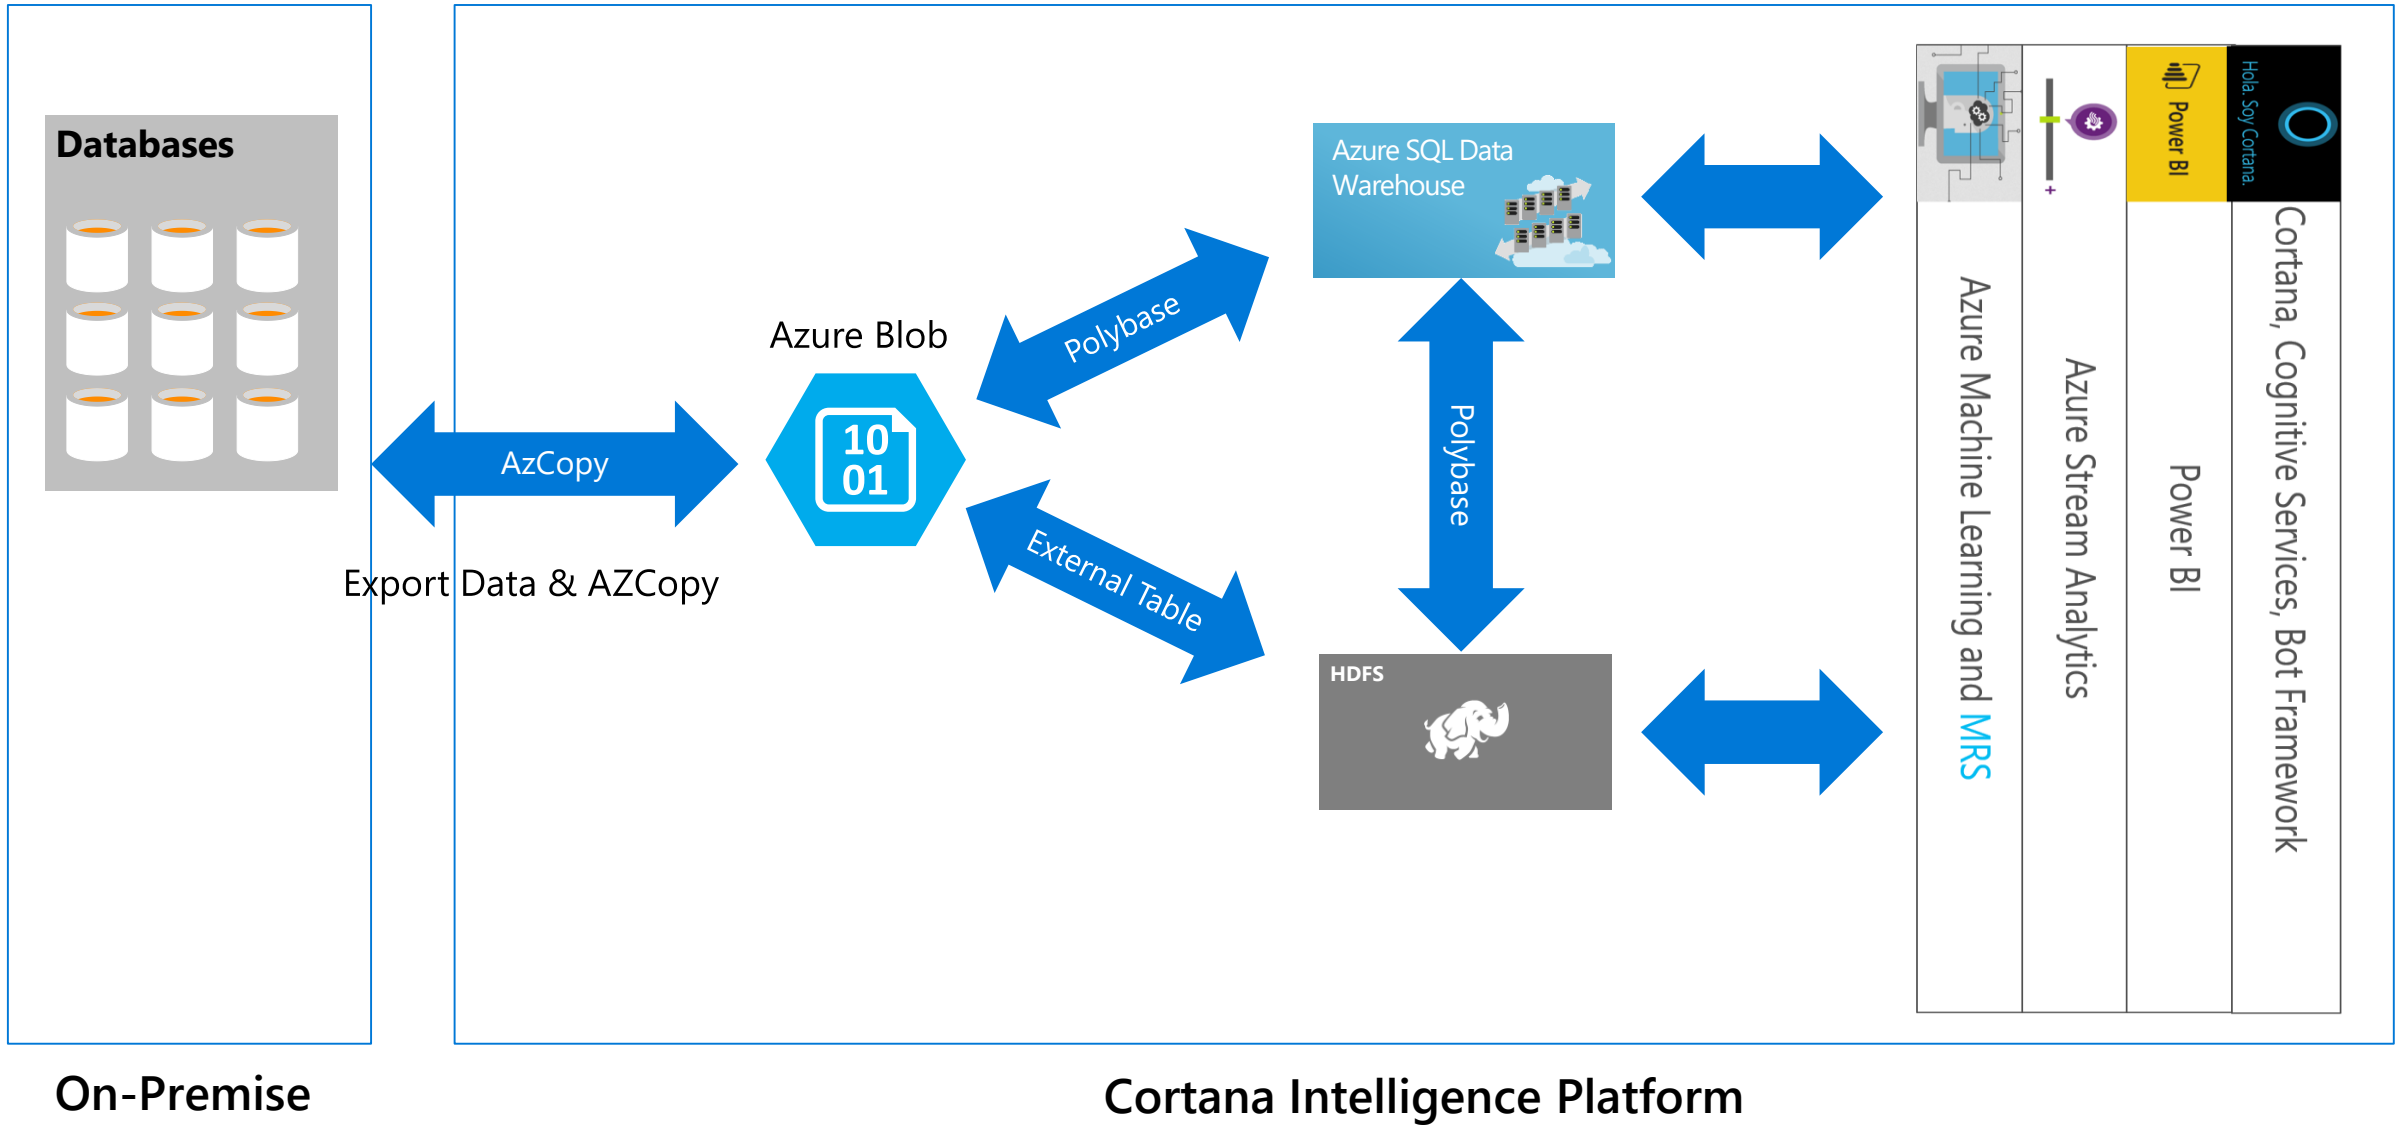

Machine Learning in ML Studio

Anomaly Detection
- One-class Support Vector Machine
- Principal Component Analysis-based Anomaly Detection
- Time Series Anomaly Detection*

Classification
- Two-class Classification
  - Averaged Perceptron
  - Bayes Point Machine
  - Boosted Decision Tree
  - Decision Forest
  - Decision Jungle
  - Logistic Regression
  - Neural Network
  - Support Vector Machine
- Multi-class Classification
  - Decision Forest
  - Decision Jungle
  - Logistic Regression
  - Neural Network
  - One-vs-all

Clustering
- K-means Clustering

Recommendation
- Matchbox Recommender

Regression
- Bayesian Linear Regression
- Boosted Decision Tree
- Decision Forest
- Fast Forest Quantile Regression
- Linear Regression
- Neural Network Regression
- Ordinal Regression
- Poisson Regression

Statistical Functions
- Descriptive Statistics
- Hypothesis Testing T-Test
- Linear Correlation
- Probability Function Evaluation

Text Analytics
- Feature Hashing
- Named Entity Recognition
- Vowpal Wabbit

Computer Vision
- OpenCV Library

https://studio.azureml.net

Guest Access Workspace: Free trial access without logging in.
Free Workspace: Free persisted access, no Azure subscription needed.
Standard Workspace: Full access with SLA under an Azure subscription.

Data/Model Visualization
- Scatterplots
- Bar Charts
- Box plots
- Histogram
- R and Python Plotting Libraries
- REPL with Jupyter Notebook
- ROC, Precision/Recall, Lift
- Confusion Matrix
- Decision Tree*

Training
- Cross Validation
- Retraining
- Parameter Sweep

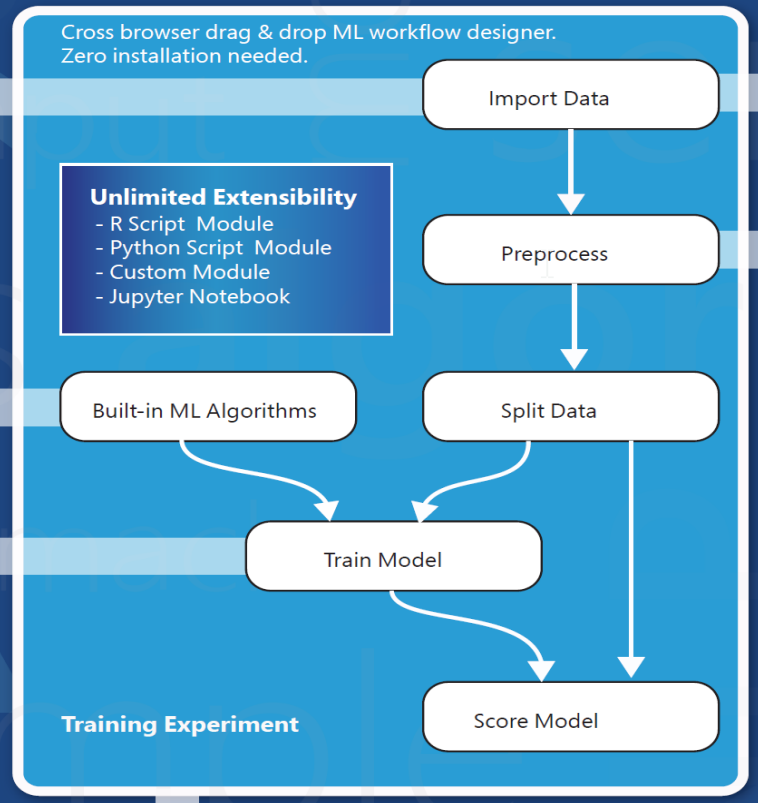

Data Source
- Azure Blob Storage
- Azure SQL DB
- Azure SQL DW*
- Azure Table
- Desktop Direct Upload
- Hadoop Hive Query
- Manual Data Entry
- OData Feed
- On-prem SQL Server*
- Web URL (HTTP)

Data Format
- ARFF
- CSV
- SVMLight
- TSV
- Excel
- ZIP

Data Preparation
- Clean Missing Data
- Clip Outliers
- Edit Metadata
- Feature Selection
- Filter
- Learning with Counts
- Normalize Data
- Partition and Sample
- Principal Component Analysis
- Quantize Data
- SQLite Transformation
- Synthetic Minority Oversampling Technique

Enterprise Grade Cloud Service
- SLA: 99.95% Guaranteed Up-time
- Azure AD Authentication
- Compute at Large Scale
- Multi-geo Availability
- Regulatory Compliance*

One-click Operationalization

Predictive Experiment

Make Prediction with Elastic APIs
- Request-Response Service (RRS)
- Batch Execution Service (BES)
- Retraining API

Community
- Gallery (http://gallery.azureml.net)
- Samples & Templates
- Workspace Sharing and Collaboration
- Live Chat & MSDN Forum Support

* Feature Coming Soon

Azure Machine Learning Studio Capabilities Overview

Created by the Azure Machine Learning Team Email: AzurePoster@microsoft.com
Download this poster: http://aka.ms/MLStudioOverview

Microsoft

Select model type based on desired algorithm

Supervised:
Make predictions based on a set of labeled examples.

Unsupervised:
No label association. Goal is to organize the data in some way or to describe its structure.

Classification:
predict a category

Regression:
a value is being predicted

Anomaly detection:
identify unusual data points

Clustering:
data segmentation

Microsoft